

Overview of a virtual cluster

using OpenNebula and SLURM

Ismael Farfán Estrada
ifarfane0900@ipn.mx

October 27, 2011

Introduction

Use virtual machines as means for “dynamic fractional resource scheduling” in a batch cluster environment.

Software used:

SLURM A resource manager, it's job is to receive batch scripts (job requests) and allocate them in the available (virtual) nodes for execution.

OpenNebula A virtual machine management service, makes it easy to work with VMs providing an easy way to create, destroy, migrate, contextualize, etc. virtual machines.

Physical connection

- The main/login server executes the SLURM resource manager and OpenNebula.
- The job of the worker nodes will be to host running VMs.
- The VMs execute the batch jobs.
- The VMs connect to the physical network using bridged connections.


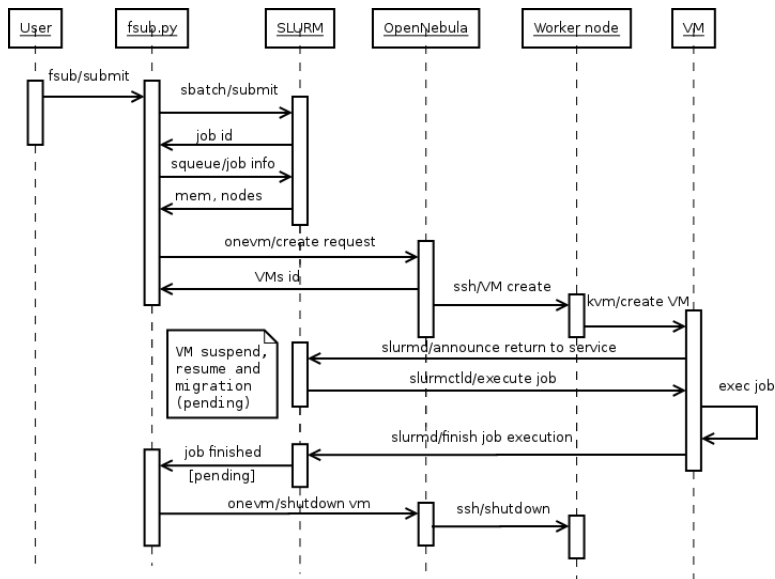


Figure: Client/server topology for the virtual cluster

Job execution sequence diagram



Job execution sequence explained

- 1 The user submits the job using `fsub.py` (Not necessary if it integrates with SLURM triggers.).
- 2 Call `sbatch` to send the job to the SLURM daemon.
- 3 Retrieve the job's requirements (nodes and memory).
- 4 Launch as many VMs as nodes requested that fit the requirements.
- 5 Wait for the VM to inform that its ready to execute the jobs.
- 6 Execute the job.
- 7 Suspend, resume, migrate VM (not working yet).
- 8 Inform job completion, cancel, kill, etc.
- 9 Shut down the VMs.

Getting job requirements

This command return the requested nodes and memory:

`squeue -o '%5D %7m' -j job-id`

```
def obten_requisitos_trabajo(tid):
    param = shlex.split( "squeue -o '%5D %7m' -j" + str(tid)
    p = subprocess.Popen(param, stdout=subprocess.PIPE)
    rt = p.communicate()
    rt = rt[0].split("\n")[1].split()
    nodos    = int( rt[0] )
    memoria  = int( rt[1] )
    if memoria < 50: memoria = 256
    return nodos, memoria
```

Launch the VMs

Create a VM template, execute the VM, register the tuple (VM-id, job-id) on the database.

```
def lanza_vms(vms, mem, tid):
    ruta = "/tmp/"
    for v in vms:
        f = open(ruta+"vm%d.one"%v, "w")
        f.write( definicion%(v, mem, v, v, v) )
        f.close()
        param = shlex.split( "onevm -v create " + f.name )
        p = subprocess.Popen(param, stdout=subprocess.PIPE)
        oid = p.communicate()
        oid = int ( oid[0].split()[1] ) # ('ID: 40\n', None)
        con.execute( "update mapa set oneid=%d, tid=%d where"
```

Cleaning upon job termination

Shutdown the VMs and update the database.

```
def limpia_trabajo(tid):  
    print "limpiando", tid  
    vms = lista_vms("where tid=%d"%tid)  
    for v in vms:  
        v = v[1]  
        print "apagando vm", v  
        param = shlex.split( "onevm -v shutdown %d"%v )  
        p = subprocess.Popen(param, stdout=subprocess.PIPE)  
        print p.communicate()  
        con.execute("update mapa set tid=-1, oneid=-1 where
```


Current test environment

- A laptop running Debian as a SLURM and OpenNebula server (with a public IP address).
- 2 worker nodes with 8-cores, 32Gb RAM and RHEL 6.0 on sierra.futuregrid.org.
- A shared directory in the NFS server of sierra for the VMs.
- Debian VMs as worker nodes.
- SQLite database for job/VM mapping purposes.

stream memory benchmark

s80r,s82r 8 core physical nodes

fg0,fg2 3 core VMs running in s80r

fg1 3 core VM running in s81r




Figure: Original and MPI stream benchmark results

What (doesn't) works

- y Create VMs
- y Job requirement retrieval
- y VM contextualization
- y VM migration
- N Live VM migration (problems with credentials, shared FS, configuration?)
- N Suspending and resuming jobs and the VMs executing them (can't setup the support, SLURM reports "node failure" (and kills the job))